

**Ex No 3                      Map Reduce program to process a weather dataset.****AIM:**

To implement MapReduce program to process a weather dataset.

**PROCEDURE:****1. Start Hadoop Services**

Make sure you are in the `sbin` folder of Hadoop. Start the Hadoop services by running the following commands:

```
cd /usr/local/Cellar/hadoop/3.4.0/libexec/sbin
./start-dfs.sh ./start-yarn.sh
```

**2. Prepare Your Files**

Create the necessary files (`mapper.py`, `reducer.py`, and `sample_weather.txt`) on your local machine or the server where Hadoop is installed.

**Create `sample_weather.txt`:**

You can create this file in the current directory:

```
nano sample_weather.txt
```

**Example data for `sample_weather.txt`:**

```
STN001 2023-09-10_04 15.0 12.0 5.0
STN001 2023-09-10_10 25.0 20.0 8.0
STN002 2023-09-10_16 30.0 25.0 10.0
STN002 2023-09-10_22 22.0 18.0 7.0
```

**Create `mapper.py`: `nano mapper.py`**

Content for `mapper.py`:

```
#!/usr/bin/python3
import sys
def map1():
```

```

for line in sys.stdin: tokens = line.strip().split() if
    len(tokens) < 4: continue station = tokens[0] date_hour =
    tokens[1] temp = tokens[2] dew = tokens[3] wind =
    tokens[4] if len(tokens) > 4 else "999.9" if temp ==
    "9999.9" or dew == "9999.9" or wind == "999.9": continue
    hour = int(date_hour.split("_")[-1]) date =
    date_hour[:date_hour.rfind("_")] if 4 < hour <= 10:
    section = "section1"
    elif 10 < hour <= 16: section
        = "section2"
    elif 16 < hour <= 22:
        section = "section3"
    else:
        section = "section4"
    key_out = f"{station}_{date}_{section}"
    value_out = f"{temp} {dew} {wind}"
    print(f"{key_out}\t{value_out}")

```

```

if __name__ == "__main__":
    map1() Create

```

**reducer.py:** nano

reducer.py

**Content for reducer.py:**

```

#!/usr/bin/python3 import
sys
def reduce1():
    current_key = None
    sum_temp, sum_dew, sum_wind = 0, 0, 0 count
    = 0

    for line in sys.stdin: key, value =
        line.strip().split("\t") temp, dew, wind =
        map(float, value.split())

        if current_key is None:

```

```

        current_key = key

    if key == current_key:
        sum_temp += temp
        sum_dew += dew
        sum_wind += wind
        count += 1
    else:
        avg_temp = sum_temp / count
        avg_dew = sum_dew / count
        avg_wind = sum_wind / count
        print(f"{current_key}\t{avg_temp} {avg_dew} {avg_wind}")

        current_key = key
        sum_temp, sum_dew, sum_wind = temp, dew, wind
        count = 1

if current_key is not None:
    avg_temp = sum_temp / count
    avg_dew = sum_dew / count
    avg_wind = sum_wind / count
    print(f"{current_key}\t{avg_temp} {avg_dew} {avg_wind}")

if __name__ == "__main__":
    reduce1()

```

### 3. Upload Files to HDFS

Next, move your data file to HDFS so that it can be processed by the Hadoop MapReduce job.

#### Create HDFS Directory:

```
hdfs dfs -mkdir /WeatherData
```

**Upload the Input Data (`sample_weather.txt`) to HDFS:** `hdfs`

```
dfs -put sample_weather.txt /WeatherData
```

#### Verify the file upload:

```
hdfs dfs -ls /WeatherData
```

You should see something like:

```
Found 1 items
-rw-r--r-- 3 user group    1234 2024-09-11 12:00
/WeatherData/sample_weather.txt
```

#### 4. Run the MapReduce Job

Now that your input file is in HDFS and your `mapper.py` and `reducer.py` are ready, you can run the MapReduce job.

Ensure you are still in the directory where your `mapper.py` and `reducer.py` scripts are located.

##### Run the Hadoop Streaming Job:

```
hadoop jar
/usr/local/Cellar/hadoop/3.4.0/libexec/share/hadoop/tools/lib/hadoop-streaming-3.4.0.jar \
-input /WeatherData/sample_weather.txt \
-output /WeatherData/output \ -mapper
"python3 mapper.py" \ reducer
"python3 reducer.py"
```

This command tells Hadoop to:

- Take the input from `/WeatherData/sample_weather.txt` on HDFS.
- Use `mapper.py` as the mapper script.
- Use `reducer.py` as the reducer script.
- Output the results to `/WeatherData/output`.

##### Note:

Ensure that both `mapper.py` and `reducer.py` have executable permissions. If not, make them executable by running:

```
chmod +x mapper.py reducer.py
```

## 5. View the Output

After the job completes, you can check the output that was stored in HDFS.

**List the output directory:**

```
hdfs dfs -ls /WeatherData/output
```

You should see something like:

```
Found 1 items
-rw-r--r--  3  user  group  456  2024-09-11  12:20
```

/WeatherData/output/part-00000 **View the output data:**

```
hdfs dfs -cat /WeatherData/output/part-00000
```

This will print the final result of the MapReduce job. You should see output similar to:

```
STN001_2023-09-10_section1 15.0 12.0 5.0
STN001_2023-09-10_section2 25.0 20.0 8.0
STN002_2023-09-10_section3 30.0 25.0 10.0
STN002_2023-09-10_section4 22.0 18.0 7.0
```

## OUTPUT:

```
nativewit@Nativewits-MacBook-Air sh$ nano sample_weather.txt
nativewit@Nativewits-MacBook-Air sh$ nano mapper.py
nativewit@Nativewits-MacBook-Air sh$ nano reducer.py
nativewit@Nativewits-MacBook-Air sh$ nano reducer.py
nativewit@Nativewits-MacBook-Air sh$ hdfs dfs -mkdir /WeatherData
2024-09-10 12:31:21,281 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
mkdir: /WeatherData: File exists
nativewit@Nativewits-MacBook-Air sh$ hdfs dfs -put sample_weather.txt /WeatherData
2024-09-10 12:31:30,969 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
put: /WeatherData/sample_weather.txt: File exists
nativewit@Nativewits-MacBook-Air sh$ hdfs dfs -ls /WeatherData
2024-09-10 12:31:43,481 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r--  1 nativewit supergroup      142  2024-09-10 10:41 /WeatherData/sample_weather.txt
nativewit@Nativewits-MacBook-Air sh$ hdfs dfs -ls /WeatherData/output
-rw-r--r--  1 nativewit supergroup      165  2024-09-10 12:32 /WeatherData/output/part-00000
nativewit@Nativewits-MacBook-Air sh$ hdfs dfs -cat /WeatherData/output/part-00000
STN001_2023-09-10_section1 15.0 12.0 5.0
STN001_2023-09-10_section2 25.0 20.0 8.0
STN002_2023-09-10_section3 30.0 25.0 10.0
STN002_2023-09-10_section4 22.0 18.0 7.0
nativewit@Nativewits-MacBook-Air sh$
```

**RESULT:**

Thus, the program for weather dataset using Map Reduce has been executed successfully.