

Ex No 8**Implement SVM/Decision tree classification techniques****AIM:**

To Implement SVM/Decision tree classification techniques using R.

PROCEDURE:

- Collect and load the dataset from sources like CSV files or databases.
- Clean and preprocess the data, including handling missing values and encoding categorical variables.
- Split the dataset into training and testing sets to evaluate model performance.
- Normalize or standardize the features, especially for SVM, to ensure consistent scaling.
- Choose the appropriate model: SVM for margin-based classification, Decision Tree for rule-based classification.
- Train the model on the training data using the 'fit' method.
- Make predictions on the testing data using the 'predict' method.
- Evaluate the model using metrics like accuracy, confusion matrix, precision, and recall.
- Visualize the results with plots, such as decision boundaries for SVM or tree structures for Decision Trees.
- Fine-tune the model by adjusting hyperparameters like 'C' for SVM or 'max_depth' for

Decision Trees.

CODE:**SVM.R:**

```
# Install and load the e1071 package (if not already installed)
install.packages("e1071") library(e1071) # Load the iris
dataset
data(iris)
# Inspect the first few rows of the dataset head(iris)
# Split the data into training (70%) and testing (30%) sets
set.seed(123) # For reproducibility sample_indices <-
sample(1:nrow(iris), 0.7 * nrow(iris)) train_data <-
iris[sample_indices, ] test_data <- iris[-sample_indices, ] # Fit the
```

```

SVM model svm_model <- svm(Species ~ ., data = train_data, kernel
= "radial") # Print the summary of the model summary(svm_model) #
Predict the test set predictions
<- predict(svm_model, newdata = test_data)
# Evaluate the model's performance confusion_matrix <- table(Predicted =
predictions, Actual = test_data$Species) print(confusion_matrix) # Calculate accuracy
accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix) cat("Accuracy:",
accuracy * 100, "%\n") Decision Tree.R:

```

```

# Install and load the rpart package (if not already installed) install.packages("rpart")
library(rpart)
# Load the iris dataset
data(iris)
# Split the data into training (70%) and testing (30%) sets
set.seed(123) # For reproducibility sample_indices <-
sample(1:nrow(iris), 0.7 * nrow(iris)) train_data <-
iris[sample_indices, ] test_data <- iris[-sample_indices, ] # Fit the
Decision Tree model tree_model <- rpart(Species ~ ., data = train_data,
method = "class")
# Print the summary of the model summary(tree_model) # Plot the
Decision Tree plot(tree_model) text(tree_model, pretty = 0) # Predict
the test set predictions <- predict(tree_model, newdata = test_data,
type = "class")
# Evaluate the model's performance confusion_matrix <- table(Predicted =
predictions, Actual = test_data$Species) print(confusion_matrix) # Calculate
accuracy accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
cat("Accuracy:", accuracy * 100, "%\n") OUTPUT: SVM in R:

```

The screenshot shows an R script in RStudio. The script loads the Iris dataset, splits it into training and testing sets, fits an SVM model with a radial kernel, and evaluates its performance. The console output shows the Fisher Scoring iterations and a confusion matrix. The Environment pane on the right lists the objects created during the process.

```

1 data(iris)
2 # Inspect the first few rows of the dataset
3 head(iris)
4 # Split the data into training (70%) and testing (30%) sets
5 set.seed(123) # For reproducibility
6 sample_indices <- sample(1:nrow(iris), 0.7 * nrow(iris))
7 train_data <- iris[sample_indices, ]
8 test_data <- iris[-sample_indices, ]
9 # Fit the SVM model
10 svm_model <- svm(Species ~ ., data = train_data, kernel = "radial")
11 # Print the summary of the model
12 summary(svm_model)
13 # Predict the test set
14 predictions <- predict(svm_model, newdata = test_data)
15 # Evaluate the model's performance
16 confusion_matrix <- table(Predicted = predictions, Actual = test_data$Species)

```

Number of Fisher Scoring iterations: 5

| | Mazda RX4 | Mazda RX4 Wag | Datsun 710 | Hornet 4 Drive |
|-------------------|------------|---------------|------------|----------------|
| | 0.46109512 | 0.46109512 | 0.59789839 | 0.49171990 |
| Hornet Sportabout | 0.29690087 | 0.25993307 | 0.09858705 | 0.70846924 |
| Merc 230 | 0.59789839 | 0.32991148 | 0.24260966 | 0.17246396 |
| Merc 450SL | 0.21552479 | 0.12601104 | 0.03197098 | 0.03197098 |
| Chrysler Imperial | 0.11005178 | 0.96591395 | 0.93878132 | 0.97821971 |
| Toyota Corona | 0.49939484 | 0.13650937 | 0.12601104 | 0.07446438 |
| Pontiac Firebird | 0.32991148 | 0.8549212 | 0.79886349 | 0.93878132 |
| Ford Pantera L | 0.14773451 | 0.36468861 | 0.11940215 | 0.49171990 |

Predicted Actual

| | setosa | versicolor | virginica |
|------------|--------|------------|-----------|
| setosa | 14 | 0 | 0 |
| versicolor | 0 | 17 | 0 |
| virginica | 0 | 1 | 13 |

Accuracy: 97.77778 %

Environment pane:

| Object | Class | Size |
|----------------|------------|-------------------------|
| data | data.frame | 7 obs. of 5 variables |
| iris | data.frame | 150 obs. of 5 variables |
| linear_model | list | List of 12 |
| logistic_model | list | List of 30 |
| mtcars | data.frame | 32 obs. of 11 variables |
| svm_model | svm | List of 31 |
| test_data | data.frame | 45 obs. of 5 variables |
| train_data | data.frame | 105 obs. of 5 variables |
| tree_model | list | List of 14 |

Decision tree:

The screenshot shows an R script in RStudio. The script loads the Iris dataset, splits it into training and testing sets, fits a Decision Tree model, and evaluates its performance. The console output shows the tree structure and a confusion matrix. The Environment pane on the right lists the objects created during the process.

```

1 Load the iris dataset
2 data(iris)
3 Split the data into training (70%) and testing (30%) sets
4 set.seed(123) # For reproducibility
5 sample_indices <- sample(1:nrow(iris), 0.7 * nrow(iris))
6 train_data <- iris[sample_indices, ]
7 test_data <- iris[-sample_indices, ]
8 # Fit the Decision Tree model
9 tree_model <- rpart(Species ~ ., data = train_data, method = "class")
10 # Print the summary of the model
11 summary(tree_model)
12 # Plot the Decision Tree
13 plot(tree_model)
14 text(tree_model, pretty = 0)
15 # Predict the test set
16 predictions <- predict(tree_model, newdata = test_data, type = "class")

```

Primary splits:

- Petal.Width < 1.75 to the left, improve=25.291950, (0 missing)
- Petal.Length < 4.75 to the left, improve=25.187810, (0 missing)
- Sepal.Length < 6.15 to the left, improve= 5.974246, (0 missing)
- Sepal.Width < 2.45 to the left, improve= 2.411006, (0 missing)

Surrogate splits:

- Petal.Length < 4.75 to the left, agree=0.913, adj=0.824, (0 split)
- Sepal.Length < 6.15 to the left, agree=0.696, adj=0.382, (0 split)
- Sepal.Width < 2.65 to the left, agree=0.638, adj=0.265, (0 split)

Node number 6: 35 observations
predicted class=versicolor expected loss=0.1142857 P(node)=0.3333333
class counts: 0 31 4
probabilities: 0.000 0.886 0.114

Node number 7: 34 observations
predicted class=virginica expected loss=0.02941176 P(node)=0.3238095
class counts: 0 1 33
probabilities: 0.000 0.029 0.971

Actual Predicted

| | setosa | versicolor | virginica |
|------------|--------|------------|-----------|
| setosa | 14 | 0 | 0 |
| versicolor | 0 | 18 | 1 |
| virginica | 0 | 0 | 12 |

Accuracy: 97.77778 %

Environment pane:

| Object | Class | Size |
|----------------|------------|-------------------------|
| data | data.frame | 7 obs. of 5 variables |
| iris | data.frame | 150 obs. of 5 variables |
| linear_model | list | List of 12 |
| logistic_model | list | List of 30 |
| mtcars | data.frame | 32 obs. of 11 variables |
| svm_model | svm | List of 31 |

RESULT:

Thus, Implement SVM and Decision tree classification techniques has been successfully executed.